

Imports & Settings

In [1]:

```
import numpy as np
import pandas as pd
import yfinance as yf

# Visualization tools
import missingno as msno
import mplfinance as mpf
import seaborn as sns
import matplotlib.pyplot as plt
%matplotlib inline

# Ignoring warnings
import warnings
warnings.filterwarnings('ignore')
```

Data Gathering

S&P 500 Constituents

In [2]:

```
snp = pd.read_csv('../data/sp500.csv')
snp.head()
```

Out[2]:

	Symbol	Security	SEC filings	GICS Sector	GICS Sub-Industry	Headquarters Location	Date first added	CIK	Founded
0	MMM	3M	reports	Industrials	Industrial Conglomerates	Saint Paul, Minnesota	1976-08-09	66740	1902
1	ABT	Abbott Laboratories	reports	Health Care	Health Care Equipment	North Chicago, Illinois	1964-03-31	1800	1888
2	ABBV	AbbVie	reports	Health Care	Pharmaceuticals	North Chicago, Illinois	2012-12-31	1551152	2013 (1888)
3	ABMD	Abiomed	reports	Health Care	Health Care Equipment	Danvers, Massachusetts	2018-05-31	815094	1981
4	ACN	Accenture	reports	Information Technology	IT Consulting & Other Services	Dublin, Ireland	2011-07-06	1467373	1989

In [3]:

```
snp_tickers = snp.Symbol.values.tolist()
```

Price Histories

In [7]:

```
prices = yf.download(snp_tickers, group_by='tickers', period='max')
```

[*****100%*****] 505 of 505 completed

1 Failed download:
- BF.B: 1d data not available for startTime=-2208988800 and endTime=1629577039. Only 100 years worth of day granularity data are allowed to be fetched per request.

It appears that the ticker BF.B failed. The '.B' portion of the ticker refers to the specific share class of the security. Referencing the share class in a ticker is not fully standardized across different financial sources and can sometimes be represented as a hyphen instead of a period. I will try to download the data for BF.B by manually changing it to BF-B instead.

In [8]:

```
bfb = yf.download(tickers='BF-B', period='max')
```

[*****100%*****] 1 of 1 completed

In [10]:

```
prices.info()
```

```
<class 'pandas.core.frame.DataFrame'>
DatetimeIndex: 15014 entries, 1962-01-02 to 2021-08-20
Columns: 3030 entries, ('COST', 'Open') to ('ZION', 'Volume')
dtypes: float64(3030)
memory usage: 347.2 MB
```

The data returned by the `yfinance.download()` function placed everything into a single dataframe with columns grouped by ticker. While useful for quickly analyzing / plotting any of the companies' histories, there are two primary issues with having everything in one dataframe:

- The file is much larger than the 100MB limit imposed by GitHub. While solutions for handling large file uploads exist, it's preferable to avoid such workarounds unless absolutely needed.
- Thinking forward to the interactive dashboard portion of this project, loading in all of the price history data for every company in the S&P 500 whenever we want to analyze just one is a waste of resources and will likely lead to unwanted load times.

By splitting the data up into individual files, both of these issues are resolved. To start, I'll export the `BF.B` / `BF-B` dataframe and then move on to looping over all of the others.

In [11]:

```
bfb.to_csv('../data/price_histories/BF-B_history.csv')
print('Successfully saved BF-B\'s history.')
```

Successfully saved BF-B's history.

In [16]:

```
for i, ticker in enumerate(snp_tickers):
    if ticker == 'BF.B':
        continue
    df = prices[ticker].dropna()
    df.to_csv(f'../data/price_histories/{ticker}_history.csv')
    print(f'{ticker.ljust(5)} -- {i+1} / {len(snp_tickers)} completed', end='\r')
```

ZTS -- 505 / 505 completed

Fundamental Data

Income Statement

In [12]:

```
df_income = pd.read_csv('../data/simfin/us-income-quarterly.csv',
                        delimiter=';',
                        parse_dates=['Report Date', 'Publish Date', 'Restated Date'])
df_income.head()
```

Out[12]:

	Ticker	SimFinId	Currency	Fiscal Year	Fiscal Period	Report Date	Publish Date	Restated Date	Shares (Basic)	Shares (Diluted)	...	Non-Operating Income (Loss)	Interest Expense, Net	Pretax Income (Loss), Adj.	Abnormal Gains (Losses)	Pretax Income (Loss)
0	A	45846	USD	2010	Q3	2010-07-31	2010-10-06	2011-09-07	347000000.0	352000000.0	...	-15000000.0	-21000000.0	100000000	127000000.0	227000000.0
1	A	45846	USD	2010	Q4	2010-10-31	2010-12-20	2012-12-20	344000000.0	356000000.0	...	35000000.0	-16000000.0	238000000	5000000.0	243000000.0
2	A	45846	USD	2011	Q1	2011-01-31	2011-03-09	2012-03-05	347000000.0	355000000.0	...	-13000000.0	-19000000.0	198000000	NaN	198000000.0
3	A	45846	USD	2011	Q2	2011-04-30	2011-06-07	2012-06-04	347000000.0	355000000.0	...	-6000000.0	-17000000.0	260000000	NaN	260000000.0
4	A	45846	USD	2011	Q3	2011-07-31	2011-09-07	2012-09-05	348000000.0	357000000.0	...	0.0	-17000000.0	281000000	NaN	281000000.0


5 rows × 28 columns



```
In [13]:
df_income.info()

<class 'pandas.core.frame.DataFrame'>
RangeIndex: 64471 entries, 0 to 64470
Data columns (total 28 columns):
#   Column                                     Non-Null Count  Dtype
---  -
0   Ticker                                     64471 non-null  object
1   SimFinId                                  64471 non-null  int64
2   Currency                                  64471 non-null  object
3   Fiscal Year                               64471 non-null  int64
4   Fiscal Period                             64471 non-null  object
5   Report Date                               64471 non-null  datetime64[ns]
6   Publish Date                              64471 non-null  datetime64[ns]
7   Restated Date                             64471 non-null  datetime64[ns]
8   Shares (Basic)                            63932 non-null  float64
9   Shares (Diluted)                          63932 non-null  float64
10  Revenue                                   63435 non-null  float64
11  Cost of Revenue                           58921 non-null  float64
12  Gross Profit                              63466 non-null  float64
13  Operating Expenses                        64313 non-null  float64
14  Selling, General & Administrative         60840 non-null  float64
15  Research & Development                    25254 non-null  float64
16  Depreciation & Amortization                29210 non-null  float64
17  Operating Income (Loss)                   64470 non-null  float64
18  Non-Operating Income (Loss)               63844 non-null  float64
19  Interest Expense, Net                     56785 non-null  float64
20  Pretax Income (Loss), Adj.                64471 non-null  int64
21  Abnormal Gains (Losses)                   34590 non-null  float64
22  Pretax Income (Loss)                      64471 non-null  int64
23  Income Tax (Expense) Benefit, Net         59898 non-null  float64
24  Income (Loss) from Continuing Operations  64471 non-null  int64
25  Net Extraordinary Gains (Losses)          10761 non-null  float64
26  Net Income                               64471 non-null  int64
27  Net Income (Common)                       64471 non-null  int64
dtypes: datetime64[ns](3), float64(15), int64(7), object(3)
memory usage: 13.8+ MB
```

```
In [25]:
msno.matrix(df_income);
```



Balance Sheet

In [10]:

```
df_balance = pd.read_csv('../data/simfin/us-balance-quarterly.csv',
                        delimiter=';',
                        parse_dates=['Report Date', 'Publish Date', 'Restated Date'])
df_balance.head()
```

Out[10]:


	Ticker	SimFinId	Currency	Fiscal Year	Fiscal Period	Report Date	Publish Date	Restated Date	Shares (Basic)	Shares (Diluted)	...	Short Term Debt	Total Current Liabilities	Long Term Debt	Total Noncurrent Liabilities	
0	A	45846	USD	2010	Q3	2010-07-31	2010-10-06	2010-10-06	347000000.0	352000000.0	...	1.501000e+09	2.917000e+09	2.177000e+09	3.373000e+09	62
1	A	45846	USD	2010	Q4	2010-10-31	2010-12-20	2011-12-16	344000000.0	356000000.0	...	1.501000e+09	3.083000e+09	2.190000e+09	3.377000e+09	64
2	A	45846	USD	2011	Q1	2011-01-31	2011-03-09	2011-03-09	347000000.0	355000000.0	...	1.000000e+06	1.406000e+09	2.138000e+09	3.299000e+09	47
3	A	45846	USD	2011	Q2	2011-04-30	2011-06-07	2011-06-07	347000000.0	355000000.0	...	0.000000e+00	1.592000e+09	2.144000e+09	3.096000e+09	46
4	A	45846	USD	2011	Q3	2011-07-31	2011-09-07	2011-09-07	348000000.0	357000000.0	...	0.000000e+00	1.505000e+09	2.168000e+09	3.048000e+09	45

5 rows × 30 columns

In [11]:

```
df_balance.info()
```

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 64471 entries, 0 to 64470
Data columns (total 30 columns):
#   Column                                                                 Non-Null Count  Dtype
---  -
0   Ticker                                                                64471 non-null object
1   SimFinId                                                             64471 non-null int64
2   Currency                                                             64471 non-null object
3   Fiscal Year                                                          64471 non-null int64
4   Fiscal Period                                                        64471 non-null object
5   Report Date                                                         64471 non-null datetime64[ns]
6   Publish Date                                                        64471 non-null datetime64[ns]
7   Restated Date                                                       64471 non-null datetime64[ns]
8   Shares (Basic)                                                      63932 non-null float64
9   Shares (Diluted)                                                    63932 non-null float64
10  Cash, Cash Equivalents & Short Term Investments                    64311 non-null float64
11  Accounts & Notes Receivable                                         59102 non-null float64
12  Inventories                                                         44891 non-null float64
13  Total Current Assets                                                64467 non-null float64
14  Property, Plant & Equipment, Net                                    63509 non-null float64
15  Long Term Investments & Receivables                                 17592 non-null float64
16  Other Long Term Assets                                              63754 non-null float64
17  Total Noncurrent Assets                                             64262 non-null float64
18  Total Assets                                                        64471 non-null int64
19  Payables & Accruals                                                 64237 non-null float64
20  Short Term Debt                                                     42547 non-null float64
21  Total Current Liabilities                                           64469 non-null float64
22  Long Term Debt                                                      51006 non-null float64
23  Total Noncurrent Liabilities                                        63706 non-null float64
24  Total Liabilities                                                   64471 non-null int64
25  Share Capital & Additional Paid-In Capital                        63379 non-null float64
26  Treasury Stock                                                      30427 non-null float64
27  Retained Earnings                                                  61579 non-null float64
28  Total Equity                                                        64470 non-null float64
29  Total Liabilities & Equity                                          64471 non-null int64
dtypes: datetime64[ns](3), float64(19), int64(5), object(3)
memory usage: 14.8+ MB
```



Cash Flow Statement

In [8]:

```
df_cashflow = pd.read_csv('../data/simfin/us-cashflow-quarterly.csv',
                           delimiter=';',
                           parse_dates=['Report Date', 'Publish Date', 'Restated Date'])
df_cashflow.head()
```

Out[8]:

	Ticker	SimFinId	Currency	Fiscal Year	Fiscal Period	Report Date	Publish Date	Restated Date	Shares (Basic)	Shares (Diluted)	...	Net Cash from Operating Activities	Change in Fixed Assets & Intangibles	Net Change in Long Term Investment	Net Cash from Acquisitions & Divestitures	Net Ca: In A
0	A	45846	USD	2010	Q3	2010-07-31	2010-10-06	2011-09-07	347000000.0	352000000.0	...	90000000.0	-27000000.0	30000000.0	-1.102000e+09	-1.1120
1	A	45846	USD	2010	Q4	2010-10-31	2010-12-20	2012-12-20	344000000.0	356000000.0	...	373000000.0	-34000000.0	0.0	-1.400000e+07	-1.4000
2	A	45846	USD	2011	Q1	2011-01-31	2011-03-09	2012-03-05	347000000.0	355000000.0	...	120000000.0	-38000000.0	5000000.0	0.000000e+00	1.5000
3	A	45846	USD	2011	Q2	2011-04-30	2011-06-07	2012-06-04	347000000.0	355000000.0	...	378000000.0	-51000000.0	9000000.0	-9.600000e+07	-1.2600
4	A	45846	USD	2011	Q3	2011-07-31	2011-09-07	2012-09-05	348000000.0	357000000.0	...	252000000.0	-32000000.0	0.0	0.000000e+00	-3.2000

5 rows × 28 columns

In [9]:


```
df_cashflow.info()

<class 'pandas.core.frame.DataFrame'>
RangeIndex: 64471 entries, 0 to 64470
Data columns (total 28 columns):
#   Column                                                                 Non-Null Count  Dtype
---  -
0   Ticker                                                                64471 non-null  object
1   SimFinId                                                             64471 non-null  int64
2   Currency                                                             64471 non-null  object
3   Fiscal Year                                                          64471 non-null  int64
4   Fiscal Period                                                        64471 non-null  object
5   Report Date                                                          64471 non-null  datetime64[ns]
6   Publish Date                                                         64471 non-null  datetime64[ns]
7   Restated Date                                                       64471 non-null  datetime64[ns]
8   Shares (Basic)                                                       63932 non-null  float64
9   Shares (Diluted)                                                     63932 non-null  float64
10  Net Income/Starting Line                                             63449 non-null  float64
11  Depreciation & Amortization                                           62356 non-null  float64
12  Non-Cash Items                                                       63349 non-null  float64
13  Change in Working Capital                                             63086 non-null  float64
14  Change in Accounts Receivable                                         5567 non-null   float64
15  Change in Inventories                                                 4126 non-null   float64
16  Change in Accounts Payable                                            4408 non-null   float64
17  Change in Other                                                       6530 non-null   float64
18  Net Cash from Operating Activities                                   64467 non-null  float64
19  Change in Fixed Assets & Intangibles                                  62995 non-null  float64
20  Net Change in Long Term Investment                                    22432 non-null  float64
21  Net Cash from Acquisitions & Divestitures                            31335 non-null  float64
22  Net Cash from Investing Activities                                    64174 non-null  float64
23  Dividends Paid                                                       32441 non-null  float64
24  Cash from (Repayment of) Debt                                         53911 non-null  float64
25  Cash from (Repurchase of) Equity                                       54287 non-null  float64
26  Net Cash from Financing Activities                                    64078 non-null  float64
27  Net Change in Cash                                                    64471 non-null  int64

dtypes: datetime64[ns](3), float64(19), int64(3), object(3)
memory usage: 13.8+ MB
```

In [31]:

```
msno.matrix(df_cashflow);
```



Merging Data

In [109]:

```
for i, ticker in enumerate(snp_tickers):
    # Manually catching 'BF.B'
    if ticker == 'BF.B':
        ticker = 'BF-B'

    # Setting dataframes
    prices = pd.read_csv(f'../data/price_histories/{ticker}_history.csv', parse_dates=['Date'])
    prices.rename(columns={'Date': 'Price Date'}, inplace=True)
    income = df_income[df_income.Ticker == ticker].set_index('Report Date')
    balance = df_balance[df_balance.Ticker == ticker].set_index('Report Date')
    cashflow = df_cashflow[df_cashflow.Ticker == ticker].set_index('Report Date')

    # Concatenating financials
    financials = pd.concat([income, balance, cashflow], axis=1).reset_index()

    # Dropping unnecessary columns
    to_drop = ['Ticker', 'SimFinId', 'Currency', 'Fiscal Year', 'Fiscal Period',
               'Publish Date', 'Restated Date', 'Shares (Basic)', 'Shares (Diluted)']
    financials = financials.drop(columns=to_drop)

    # Merging in prices
    final_df = pd.merge_asof(left=financials,
                             right=prices,
                             left_on='Report Date',
                             right_on='Price Date',
                             direction='backward')

    # Exporting dataframe
    final_df.to_csv(f'../data/merged_data/{ticker}_merged.csv', index=False)

    # Printing progress
    print(f'{i+1}/{len(snp_tickers)} -- merged and saved {ticker.ljust(5)}', end='\r')
```

505/505 -- merged and saved ZTS

Loading in an example to ensure everything worked properly:

In [112]:

```
msft_merged = pd.read_csv('../data/merged_data/MSFT_merged.csv', parse_dates=['Report Date', 'Price Date'])
msft_merged
```

Out[112]:

	Report Date	Revenue	Cost of Revenue	Gross Profit	Operating Expenses	Selling, General & Administrative	Research & Development	Depreciation & Amortization	Operating Income (Loss)	Non-Operating Income (Loss)	...	Cash from (Repurchase of) Equity
0	2004-03-31	9.175000e+09	-1.411000e+09	7.764000e+09	-6.486000e+09	-4.948000e+09	-1.538000e+09	NaN	1.278000e+09	1.001000e+09	...	-1.040000e+09
1	2004-06-30	9.292000e+09	-1.361000e+09	7.931000e+09	-4.798000e+09	-3.183000e+09	-1.615000e+09	NaN	3.133000e+09	5.710000e+08	...	1.296000e+09
2	2004-09-30	9.189000e+09	-1.405000e+09	7.784000e+09	-4.290000e+09	-2.760000e+09	-1.530000e+09	NaN	3.494000e+09	2.790000e+08	...	1.320000e+08
3	2004-12-31	1.081800e+10	-1.875000e+09	8.943000e+09	-4.194000e+09	-2.773000e+09	-1.421000e+09	NaN	4.749000e+09	4.200000e+08	...	-1.740000e+08
4	2005-03-31	9.620000e+09	-1.363000e+09	8.257000e+09	-4.928000e+09	-3.446000e+09	-1.482000e+09	NaN	3.329000e+09	4.960000e+08	...	-2.073000e+09
...
61	2019-06-30	3.371700e+10	-1.041200e+10	2.330500e+10	-1.090000e+10	-6.387000e+09	-4.513000e+09	NaN	1.240500e+10	1.910000e+08	...	-4.325000e+09
62	2019-09-30	3.305500e+10	-1.040600e+10	2.264900e+10	-9.963000e+09	-5.398000e+09	-4.565000e+09	NaN	1.268600e+10	0.000000e+00	...	-4.485000e+09
63	2019-12-31	3.690600e+10	-1.235800e+10	2.454800e+10	-1.065700e+10	-6.054000e+09	-4.603000e+09	NaN	1.389100e+10	1.940000e+08	...	-4.972000e+09
64	2020-03-31	3.502100e+10	-1.097500e+10	2.404600e+10	-1.107100e+10	-6.184000e+09	-4.887000e+09	NaN	1.297500e+10	-1.320000e+08	...	-6.717000e+09
65	2020-06-30	3.803300e+10	-1.233900e+10	2.569400e+10	-1.228700e+10	-7.073000e+09	-5.214000e+09	NaN	1.340700e+10	1.500000e+07	...	-5.451000e+09

66 rows × 64 columns



```
In [113]:  
msft_merged.info()
```

```
<class 'pandas.core.frame.DataFrame'>  
RangeIndex: 66 entries, 0 to 65  
Data columns (total 64 columns):  
#   Column                                     Non-Null Count  Dtype  
---  -  
0   Report Date                             66 non-null    datetime64[ns]  
1   Revenue                                66 non-null    float64  
2   Cost of Revenue                        66 non-null    float64  
3   Gross Profit                           66 non-null    float64  
4   Operating Expenses                     66 non-null    float64  
5   Selling, General & Administrative      66 non-null    float64  
6   Research & Development                  66 non-null    float64  
7   Depreciation & Amortization             0 non-null     float64  
8   Operating Income (Loss)                 66 non-null    float64  
9   Non-Operating Income (Loss)             66 non-null    float64  
10  Interest Expense, Net                   0 non-null     float64  
11  Pretax Income (Loss), Adj.              66 non-null    int64  
12  Abnormal Gains (Losses)                 17 non-null    float64  
13  Pretax Income (Loss)                    66 non-null    int64  
14  Income Tax (Expense) Benefit, Net       66 non-null    float64  
15  Income (Loss) from Continuing Operations 66 non-null    int64  
16  Net Extraordinary Gains (Losses)        0 non-null     float64  
17  Net Income                              66 non-null    int64  
18  Net Income (Common)                     66 non-null    int64  
19  Cash, Cash Equivalents & Short Term Investments 66 non-null    float64  
20  Accounts & Notes Receivable             66 non-null    float64  
21  Inventories                             66 non-null    float64  
22  Total Current Assets                     66 non-null    float64  
23  Property, Plant & Equipment, Net        66 non-null    float64  
24  Long Term Investments & Receivables      66 non-null    float64  
25  Other Long Term Assets                  66 non-null    float64  
26  Total Noncurrent Assets                  66 non-null    float64  
27  Total Assets                             66 non-null    int64  
28  Payables & Accruals                     66 non-null    float64  
29  Short Term Debt                         45 non-null    float64  
30  Total Current Liabilities                66 non-null    float64  
31  Long Term Debt                          45 non-null    float64  
32  Total Noncurrent Liabilities            66 non-null    float64  
33  Total Liabilities                       66 non-null    int64  
34  Share Capital & Additional Paid-In Capital 66 non-null    float64  
35  Treasury Stock                          0 non-null     float64  
36  Retained Earnings                       66 non-null    float64  
37  Total Equity                             66 non-null    float64  
38  Total Liabilities & Equity               66 non-null    int64  
39  Net Income/Starting Line                 66 non-null    float64  
40  Depreciation & Amortization.1           66 non-null    float64  
41  Non-Cash Items                          66 non-null    float64  
42  Change in Working Capital                66 non-null    float64  
43  Change in Accounts Receivable            66 non-null    float64  
44  Change in Inventories                    45 non-null    float64  
45  Change in Accounts Payable              45 non-null    float64  
46  Change in Other                         66 non-null    float64  
47  Net Cash from Operating Activities        66 non-null    float64  
48  Change in Fixed Assets & Intangibles     66 non-null    float64  
49  Net Change in Long Term Investment       66 non-null    float64  
50  Net Cash from Acquisitions & Divestitures 64 non-null    float64  
51  Net Cash from Investing Activities        66 non-null    float64  
52  Dividends Paid                           65 non-null    float64  
53  Cash from (Repayment of) Debt            48 non-null    float64  
54  Cash from (Repurchase of) Equity          66 non-null    float64  
55  Net Cash from Financing Activities        66 non-null    float64  
56  Net Change in Cash                       66 non-null    int64  
57  Price Date                              66 non-null    datetime64[ns]  
58  Open                                    66 non-null    float64  
59  High                                    66 non-null    float64  
60  Low                                    66 non-null    float64  
61  Close                                   66 non-null    float64  
62  Adj Close                              66 non-null    float64  
63  Volume                                  66 non-null    float64  
dtypes: datetime64[ns](2), float64(53), int64(9)  
memory usage: 33.1 KB
```

Exploratory Data Analysis

Plotting Price Histories

In [184]:

```
def plot_price_history(ticker, volume=False):
    prices = pd.read_csv(f'../data/price_histories/{ticker}_history.csv',
                        index_col='Date',
                        parse_dates=True)
    fig, ax = mpf.plot(prices,
                      type='line',
                      volume=volume,
                      returnfig=True,
                      figscale=0.75)
    ax[0].set_title(f'{ticker} Price History')
```

In [185]:

```
plot_price_history('GOOGL')
```



In [186]:

```
plot_price_history('MSFT')
```



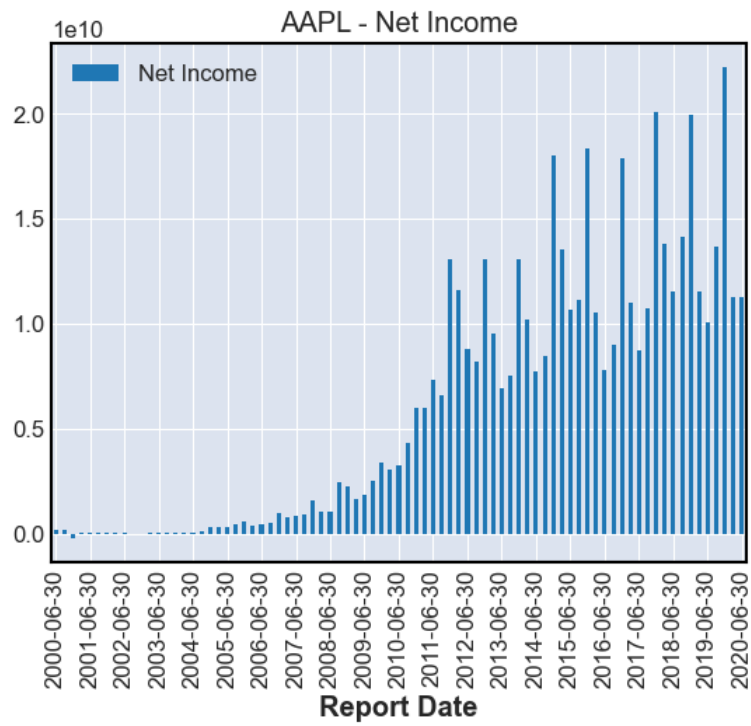
Plotting Fundamental Histories

In [236]:

```
def plot_fundamental_history(ticker, fundamental):
    df = pd.read_csv(f'../data/merged_data/{ticker}_merged.csv', index_col='Report Date')
    df = df[[fundamental]]
    df.plot(kind='bar')
    plt.xticks(ticks=np.arange(0, len(df), 4), labels=df.index[:4])
    plt.title(f'{ticker} - {fundamental}')
```

In [237]:

```
plot_fundamental_history('AAPL', 'Net Income')
```




Scatter Plots of Price vs Fundamental

In [153]:

```
def plot_scatter(ticker, fundamental):  
    df = pd.read_csv(f'../data/merged_data/{ticker}_merged.csv')  
    fundamental = df[fundamental]  
    price = df.Close  
    sns.scatterplot(fundamental, price)  
    plt.title(ticker)
```


In [187]:

```
plot_scatter('XYL', 'Revenue')
```



In [188]:

```
plot_scatter('MSFT', 'Research & Development')
```



Data Preprocessing

Handling Missing Values

The primary preprocessing need is handling missing values. Imputation is a somewhat risky methodology for this particular data, even with more advanced techniques such as iterative imputation. While not ideal, the safest route is to drop those columns which do not meet a certain minimum non-null threshold and then the remaining rows with null values thereafter. This threshold will be set at 30 in order to ensure there is enough data available for the model to work with. This will likely exclude some companies from being able to be modeled but this is acceptable for a proof of concept.

In [48]:

```
for i, ticker in enumerate(snp_tickers):
    # Manually catching 'BF.B'
    if ticker == 'BF.B':
        ticker = 'BF-B'

    # Loading in data
    df = pd.read_csv(f'../data/merged_data/{ticker}_merged.csv')

    # Dropping columns below the threshold
    threshold = 30
    to_drop = [col for col in df.columns if df[col].isna().sum() > len(df) - threshold]
    df = df.drop(columns=to_drop)

    # Dropping rows with remaining missing data
    df = df.dropna()

    # Exporting dataframe
    df.to_csv(f'../data/preprocessed_data/{ticker}_preprocessed.csv', index=False)

    # Printing progress
    print(f'{i+1}/{len(snp_tickers)} -- processed and saved {ticker.ljust(5)}', end='\r')
```

505/505 -- processed and saved ZTS

Number of Companies Excluded by Threshold Value

In [49]:

```
from pandas.errors import EmptyDataError
```

In [50]:

```
empty = []
for ticker in snp_tickers:
    if ticker == 'BF.B':
        ticker = 'BF-B'
    try:
        df = pd.read_csv(f'../data/preprocessed_data/{ticker}_preprocessed.csv')
    except EmptyDataError as e:
        empty.append(ticker)
```

In [51]:

```
len(empty)
```

Out[51]:

151